


SODIFRANCE/NETAPSYS VOUS SOUHAITE LA BIENVENUE

Spécialiste de la
Transformation des SI



1300

COLLABORATEURS

16

IMPLANTATIONS



+150

Projets réalisés

Précurseur des
architectures micros
services



+100

M€ d'investissements
en R&D
depuis 20 ans



Plusieurs milliards de
lignes de code traitées

Des solutions présentes dans

14 pays



Le traitement des mesures des applications cloud-natives



MOST METRICS ARE BETTER THOUGHT OF AS DISTRIBUTIONS RATHER THAN AVERAGES

“ Monitoring and alerting based only on the average latency would show no change in behavior ”




Figure 4-1. 50th, 85th, 95th, and 99th percentile latencies for a system. Note that the Y-axis has a logarithmic scale.

“ Using percentiles for indicators allows you to consider the shape of the distribution and its differing attributes: a high-order percentile, such as the **99th or 99.9th**, shows you a **plausible worst-case value**, while using the 50th percentile (also known as the median) emphasizes the typical case. **The higher the variance in response times, the more the typical user experience is affected by long-tail behavior**, an effect exacerbated at high load by **queuing effects**. User studies have shown that people typically prefer a slightly slower system to one with high variance in response time, so some **SRE teams focus only on high percentile values, on the grounds that if the 99.9th percentile behavior is good, then the typical experience is certainly going to be.** ”

ROB PIKE'S 5 RULES OF PROGRAMMING

- **Rule 1.** You can't tell where a program is going to spend its time. Bottlenecks occur in surprising places, so don't try to second guess and put in a speed hack until you've proven that's where the bottleneck is.
- **Rule 2. Measure.** Don't tune for speed until you've measured, and even then don't unless one part of the code overwhelms the rest.
- **Rule 3.** Fancy algorithms are slow when n is small, and n is usually small. Fancy algorithms have big constants. Until you know that n is frequently going to be big, don't get fancy. (Even if n does get big, use Rule 2 first.)
- **Rule 4.** Fancy algorithms are buggier than simple ones, and they're much harder to implement. Use simple algorithms as well as simple data structures.
- **Rule 5.** Data dominates. If you've chosen the right data structures and organized things well, the algorithms will almost always be self-evident. Data structures, not algorithms, are central to programming.

Pike's rules 1 and 2 restate Tony Hoare's famous maxim "Premature optimization is the root of all evil." Ken Thompson rephrased Pike's rules 3 and 4 as "When in doubt, use brute force.". Rules 3 and 4 are instances of the design philosophy KISS. Rule 5 was previously stated by Fred Brooks in The Mythical Man-Month. Rule 5 is often shortened to "write stupid code that uses smart objects".

SYNOPTIQUE DU SIMULATEUR

Ip:port/metrics

Injection des paramètres pour simuler une rampe ascendante et descendante



Initialisation d'un sémaphore avec le nombre maximum de goroutines simultanées

Sémaphore → true
WG ++
go bubble(sem)

Durée du tir

Attente de la fin des goroutines (WG)

Paramétrage du comportement par ligne de commande (nb de cœurs, pente, durée, nb de goroutines, nom, ...)

Bubble()

Trie à bulle + wait (goroutine)

Sémaphore ←
WG --

Mesure

Mesure

Mesure

Mesure

Mesure

Mesure

Mesure

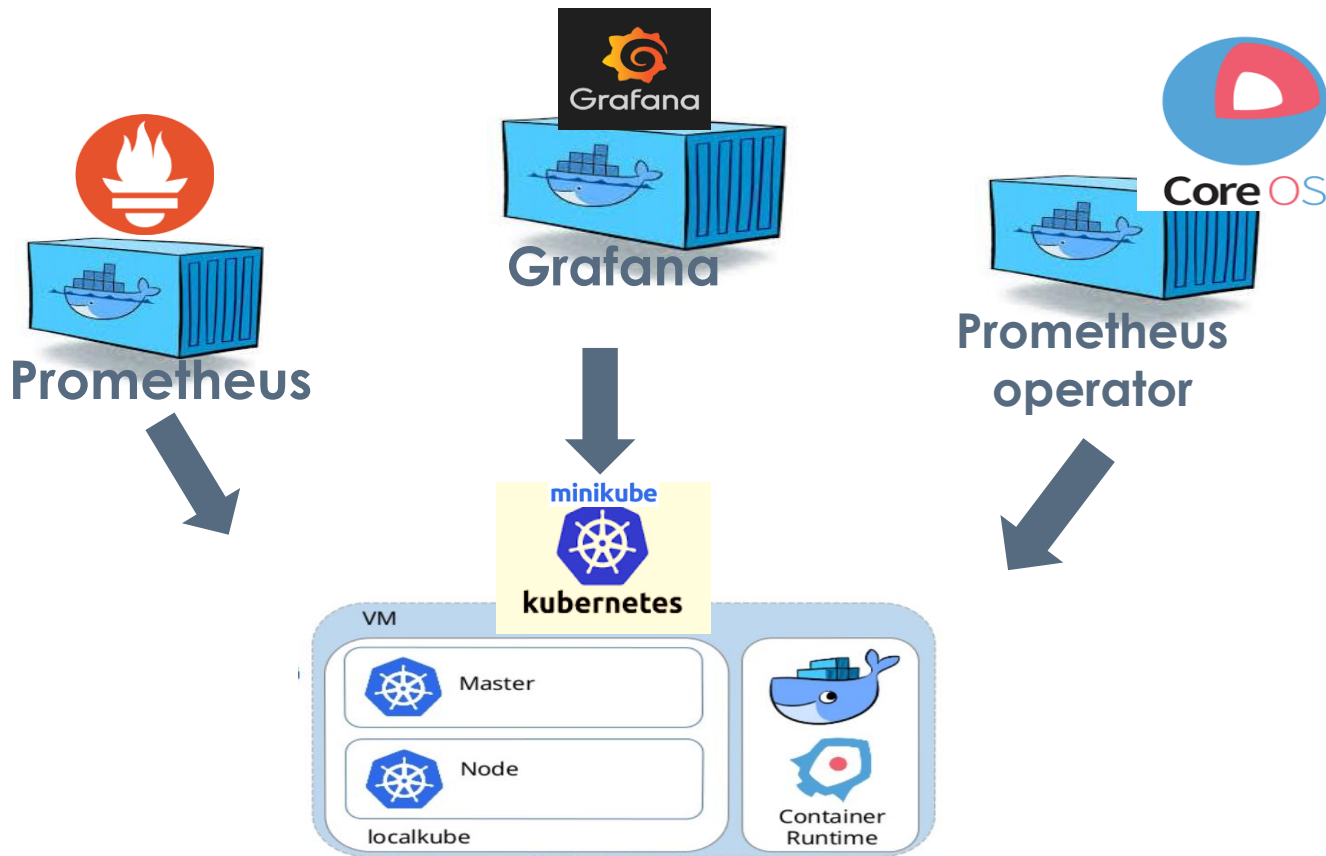
Mesure

Expose

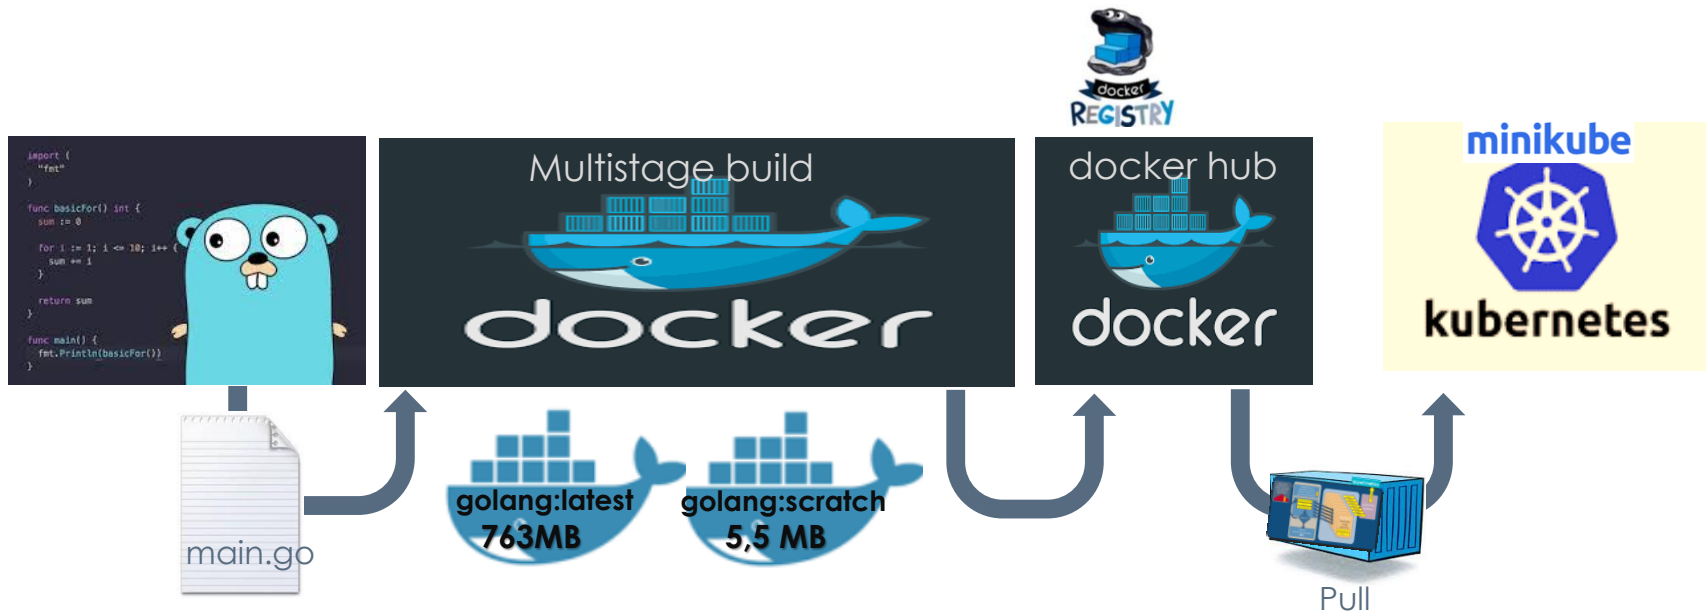
goroutine grafana

DEMO

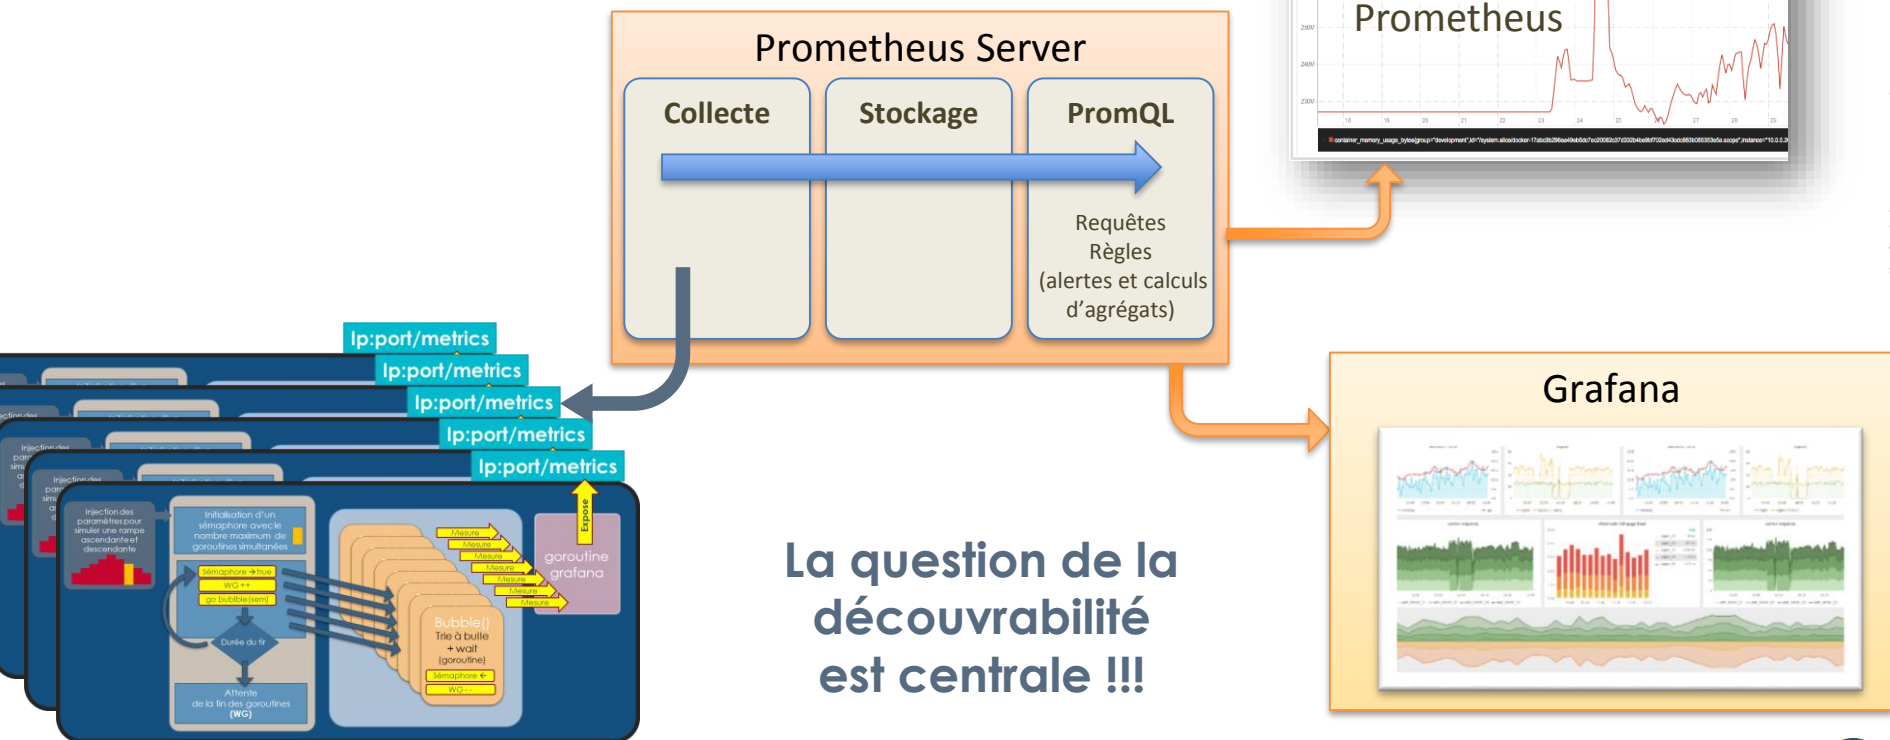
BOOTSTRAP DU CLUSTER




DU CODE AU DÉPLOIEMENT DANS LE CLUSTER



DU CLUSTER AU TABLEAU DE BORD







Opérations par seconde




Consommation CPU




Injecteur nombre de goroutine en //




Allocation RAM



Temps de traitement



GC



CONCLUSION

- Le traitement des mesures est une chose ...
- Nous n'avons travaillé que sur la détectabilité des problèmes
- Maintenant, il faut les analyser pour les résoudre
- C'est le job des traces
 - ▶ Comme pour les mesures, elles doivent être adaptées aux applications cloud native
 - ▶ <https://www.cncf.io/projects/>



OpenTracing
Distributed Tracing API



Jaeger
Distributed Tracing



distributed transaction
monitoring



performance and
latency optimization



root cause analysis



service dependency
analysis



distributed context
propagation



<https://github.com/jaegertracing/jaeger/tree/master/examples/hotrod>

QUELQUES LIENS ...

- <https://github.com/marcdivet/Bubble>
- <https://prometheus.io/docs/introduction/overview/>
- <https://grafana.com/>
- <https://coreos.com/blog/the-prometheus-operator.html>
 - ▶ <https://coreos.com/operators/prometheus/docs/latest/user-guides/getting-started.html>
 - ▶ <https://coreos.com/operators/prometheus/docs/latest/>
- https://github.com/prometheus/node_exporter
- <https://github.com/kubernetes/kube-state-metrics/>
- <https://kubernetes.io/docs/getting-started-guides/minikube/>
- <https://docs.docker.com/engine/installation/>
- <http://blog.cloud66.com/how-to-create-the-smallest-possible-docker-image-for-your-golang-application/>
 - ▶ <https://www.habitus.io>

C'EST FINI ... MAINTENANT :

